

## **Program 1**

Implement Naïve Bayes models and Bayesian networks. (Demonstrate the diagnosis of heart patients using standard heart disease data set etc)

### **Algorithm:**

- Load and preprocess the dataset (e.g., UCI Heart Disease dataset).
- Calculate prior probabilities for each class (disease present or not).
- Compute likelihood probabilities using conditional probability and the assumption of feature independence (for Naïve Bayes).
- Apply Bayes' Theorem to compute posterior probabilities.
- Classify new patient data based on the highest posterior probability

### **Code:**

```
import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.naive_bayes import GaussianNB

from sklearn.metrics import accuracy_score, classification_report

from pgmpy.models import BayesianModel

from pgmpy.estimators import MaximumLikelihoodEstimator

from pgmpy.inference import VariableElimination

dataset_url = "https://archive.ics.uci.edu/ml/machine-learning-databases/heart-disease/processed.cleveland.data"

columns = ["age", "sex", "cp", "trestbps", "chol", "fbs", "restecg", "thalach", "exang", "oldpeak", "slope", "ca", "thal", "target"]

data = pd.read_csv(dataset_url, names=columns, na_values='?')

data.dropna(inplace=True)

data["target"] = (data["target"] > 0).astype(int)

X = data.drop(columns=["target"])
```

```

y = data["target"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

nb_model = GaussianNB()

nb_model.fit(X_train, y_train)

y_pred = nb_model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))

print(classification_report(y_test, y_pred))

model = BayesianModel([("age", "target"), ("sex", "target"), ("cp", "target"), ("chol", "target"), ("thal", "target")])

model.fit(data, estimator=MaximumLikelihoodEstimator)

infer = VariableElimination(model)

print("Probability of Heart Disease given cp=3:")

print(infer.query(variables=["target"], evidence={"cp": 3}))

```

## Output

```

Accuracy: 0.9166666666666666
              precision    recall  f1-score   support

     0           0.90      0.97      0.93         36
     1           0.95      0.83      0.89         24

 accuracy                   0.92         60
 macro avg              0.92      0.90      0.91         60
 weighted avg           0.92      0.92      0.92         60

```

## **Program 2**

Implement a simple linear regression algorithm to predict a continuous target variable based on a given dataset.

### **Algorithm:**

- Load dataset and preprocess it.
- Define the hypothesis function:  $Y=mX+b$ .
- Compute the cost function (Mean Squared Error).
- Use Gradient Descent or Least Squares to optimize  $m$  and  $b$ .
- Predict values for new inputs and visualize results.

### **Code:**

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score

from sklearn.datasets import fetch_california_housing

data = fetch_california_housing()

df = pd.DataFrame(data.data, columns=data.feature_names)

df["Price"] = data.target

X = df[["MedInc"]]

y = df["Price"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LinearRegression()
```

```
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared: {r2:.2f}")
```

## Output

```
Mean Squared Error: 0.71
R-squared: 0.46
```

### **Program 3**

Develop a program to implement a Support Vector Machine for binary classification. Use a sample dataset and visualize the decision boundary.

#### **Algorithm:**

- Load and preprocess dataset.
- Define the SVM objective: maximize margin between two classes.
- Solve the optimization problem using techniques like the SMO algorithm.
- Use kernels (linear, polynomial, RBF) if needed.
- Plot the decision boundary using Matplotlib.

#### **Code:**

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn import datasets

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score

from mlxtend.plotting import plot_decision_regions

iris = datasets.load_iris()

df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

df['target'] = iris.target

df = df[df['target'] != 2]

X = df[['sepal length (cm)', 'sepal width (cm)']].values

y = df['target'].values
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

svm_model = SVC(kernel='linear')

svm_model.fit(X_train, y_train)

y_pred = svm_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')

plt.figure(figsize=(8,6))

plot_decision_regions(X_train, y_train, clf=svm_model)

plt.xlabel('Sepal Length (cm)')


plt.ylabel('Sepal Width (cm)')

plt.title('SVM Decision Boundary (Iris Dataset)')

plt.show()

```

## Output



## **Program 4**

Write a program to demonstrate the ID3 decision tree algorithm using an appropriate dataset for classification.

### **Algorithm:**

- Load dataset and preprocess it.
- Compute entropy and information gain for each feature.
- Select the feature with the highest information gain as the root node.
- Recursively split the dataset based on feature values until all instances belong to the same class.
- Use the trained tree to classify new data points.

### **Code:**

```
import pandas as pd

from google.colab import drive

drive.mount('/content/drive')

data=pd.read_csv('/content/drive/MyDrive/workshop/placementdata.csv')

import numpy as np

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.tree import DecisionTreeClassifier, plot_tree

from sklearn.metrics import accuracy_score,confusion_matrix,classification_report

data.head()

data.duplicated().sum()

data.isnull().sum()

data.shape
```

```

data.info()

data.describe()

data['PlacementStatus'].value_counts().plot(kind='bar')

data['PlacementStatus'].value_counts()/len(data)*100

placement_plot=['Internships','Projects','Workshops/Certifications','ExtracurricularActivities','Placement
Training']

plt.figure(figsize=(15,12))

for i,graph in enumerate(placement_plot):

    plt.subplot(2,3,i+1)

    sns.countplot(x=graph,data=data,hue='PlacementStatus')

    plt.title(f'Placement status based on {graph}')

num_col=data.select_dtypes('number')

num_col.columns.tolist()

non_num_col=data.select_dtypes('object')

non_num_col.columns.tolist()

data['PlacementStatus'].replace({'Placed':1,'NotPlaced':0},inplace=True)

data['PlacementStatus'].head()

from sklearn.preprocessing import LabelEncoder

le=LabelEncoder()

data['PlacementStatus']=le.fit_transform(data['PlacementStatus'])

data['ExtracurricularActivities']=le.fit_transform(data['ExtracurricularActivities'])

data['PlacementTraining']=le.fit_transform(data['PlacementTraining'])

data.info()

x=data.drop(['PlacementStatus','StudentID'],axis=1)

y=data['PlacementStatus']

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=42)

```



```

scaler=StandardScaler()

scaler=StandardScaler()

x_train_scaled=scaler.fit_transform(x_train)

x_test_scaled=scaler.transform(x_test)

model=DecisionTreeClassifier()

model.fit(x_train_scaled,y_train)

plt.figure(figsize=(12,8))

plot_tree(model,filled=True,feature_names=x.columns,class_names=['Placed','NotPlaced'])

plt.show()

clf=DecisionTreeClassifier(max_depth=3,min_samples_leaf=5,min_samples_split=10)

clf.fit(x_train_scaled,y_train)

plt.figure(figsize=(12,8))

plot_tree(clf,filled=True,feature_names=x.columns,class_names=['Placed','NotPlaced'])

plt.show()

y_pred_model=model.predict(x_test_scaled)

y_pred_clf=clf.predict(x_test_scaled)

print(classification_report(y_test,y_pred_model))

print(classification_report(y_test,y_pred_clf))

importances=model.feature_importances_

feature_importances_df=pd.DataFrame({'feature':x_train.columns,'importance':importances})

feature_importances_df.sort_values(by='importance',ascending=False,inplace=True)

feature_importances_df

x=data.drop(['SoftSkillsRating','Workshops/Certifications','Internships','Projects','PlacementTraining'],axis=1)

print(x.head())

x=x.drop(['SoftSkillsRating','SoftSkillsRating','Workshops/Certifications','Internships','Projects','Placem

```

```

entTraining'],axis=1)

y=data['PlacementStatus']

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=42)

scaler=StandardScaler()

scaler=StandardScaler()

x_train_scaled=scaler.fit_transform(x_train)

x_test_scaled=scaler.transform(x_test)

clf=DecisionTreeClassifier(max_depth=3,min_samples_leaf=5,min_samples_split=10)

clf.fit(x_train_scaled,y_train)

plt.figure(figsize=(12,8))

plot_tree(clf,filled=True,feature_names=x.columns,class_names=['Placed','NotPlaced'])


plt.show()

y_pred_clf=clf.predict(x_test_scaled)

print(classification_report(y_test,y_pred_clf))

```

# Output



## **Program 5**

Implement a KNN algorithm for regression tasks instead of classification. Use a small dataset, and predict continuous values based on the average of the nearest neighbors.

### **Algorithm:**

- Load and preprocess dataset.
- Choose the number of neighbors (K).
- Compute the Euclidean distance between the query point and all training samples.
- Select the K nearest neighbors.
- Compute the average of their target values to predict the output.

### **Code:**

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsRegressor

from sklearn.metrics import mean_squared_error, r2_score

from sklearn.datasets import fetch_california_housing

data = fetch_california_housing()

df = pd.DataFrame(data.data, columns=data.feature_names)

df["Price"] = data.target # Target variable

X = df[["MedInc"]] # Median income as predictor

y = df["Price"] # House price as target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

knn_model = KNeighborsRegressor(n_neighbors=5)

knn_model.fit(X_train, y_train)
```

```

y_pred = knn_model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)

r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse:.2f} ')

print(f'R-squared: {r2:.2f} ')

plt.scatter(X_test, y_test, color='blue', label='Actual data')

plt.scatter(X_test, y_pred, color='red', label='Predicted values', alpha=0.6)

plt.xlabel("Median Income")

plt.ylabel("House Price")

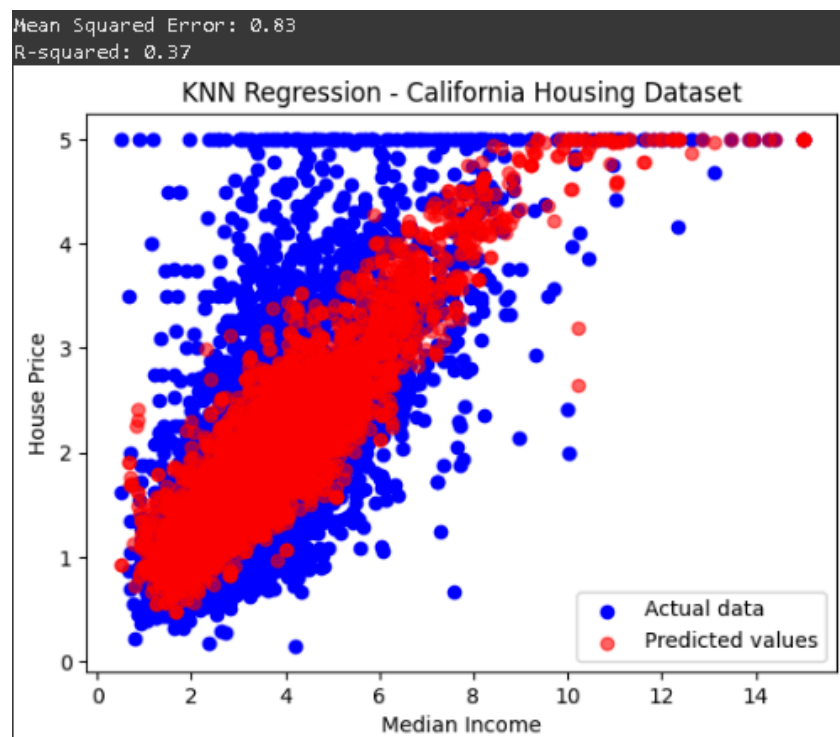
plt.title("KNN Regression - California Housing Dataset")

plt.legend()

plt.show()

```

## Output



## **Program 6**

Implement the k-Nearest Neighbor algorithm to classify the Iris dataset, printing both correct and incorrect predictions

### **Algorithm:**

- Load the Iris dataset.
- Choose K.
- Compute the Euclidean distance between the query point and training points.
- Select K nearest neighbors.
- Assign the most common class among the neighbors to the query point.
- Print correct and incorrect predictions.

### **Code:**

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn import datasets

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score, classification_report

iris = datasets.load_iris()

df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

df['target'] = iris.target

X = df[iris.feature_names].values

y = df['target'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

knn_model = KNeighborsClassifier(n_neighbors=5)
```

```

knn_model.fit(X_train, y_train)

y_pred = knn_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')

print(classification_report(y_test, y_pred))

```

## Output

```

Accuracy: 1.00
              precision    recall  f1-score   support

     0           1.00        1.00        1.00         10
     1           1.00        1.00        1.00          9
     2           1.00        1.00        1.00         11

   accuracy                   1.00          30
  macro avg           1.00        1.00        1.00          30
 weighted avg           1.00        1.00        1.00          30

```

## **Program 7**

Develop a program to implement the non-parametric Locally Weighted Regression algorithm, fitting data points and visualizing results.

### **Algorithm:**

- Load and preprocess dataset.
- Define a weight function  $W(x)$  that assigns higher weights to closer points.
- Compute the weighted least squares estimate for a given query point.
- Predict and visualize the regression line.

### **Code:**

```
import numpy as np

import matplotlib.pyplot as plt

def gaussian_kernel(x, x0, tau):

    return np.exp(-np.sum((x - x0)**2) / (2 * tau**2))

def compute_weights(X, x0, tau):

    m = X.shape[0]

    weights = np.zeros(m)

    for i in range(m):

        weights[i] = gaussian_kernel(X[i], x0, tau)

    return np.diag(weights)

def locally_weighted_regression(X, y, x0, tau):

    X_b = np.c_[np.ones((X.shape[0], 1)), X] # Add intercept term

    x0_b = np.r_[1, x0] # Add intercept term to the query point

    W = compute_weights(X, x0, tau)

    theta = np.linalg.inv(X_b.T @ W @ X_b) @ (X_b.T @ W @ y)

    return x0_b @ theta
```



```

def plot_lwr(X, y, tau):

    X_range = np.linspace(np.min(X), np.max(X), 300)

    y_pred = [locally_weighted_regression(X, y, x0, tau) for x0 in X_range]

    plt.scatter(X, y, color='blue', label='Data points')

    plt.plot(X_range, y_pred, color='red', label='LWR fit')

    plt.xlabel('X')

    plt.ylabel('y')

    plt.title(f'Locally Weighted Regression (tau={tau})')

    plt.legend()

    plt.show()


X = np.array([[1], [2], [3], [4], [5]])

y = np.array([1, 3, 2, 5, 4])

plot_lwr(X, y, tau=0.5)

```

## Output



## **Program 8**

Build an Artificial Neural Network by implementing the Backpropagation algorithm and test it with suitable datasets.

### **Algorithm:**

- Initialize weights randomly.
- Forward propagation: Compute activations using weights.
- Compute error using a loss function (e.g., Mean Squared Error for regression, Cross-Entropy for classification).
- Backpropagation: Compute gradients and update weights using Gradient Descent.
- Repeat until convergence.

### **Code:**

```
import numpy as np

import matplotlib.pyplot as plt

from sklearn.datasets import make_moons

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import OneHotEncoder

X, y = make_moons(n_samples=500, noise=0.2, random_state=42)

y = y.reshape(-1, 1)

encoder = OneHotEncoder(sparse_output=False)

y = encoder.fit_transform(y)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

def sigmoid(x):

    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):

    return x * (1 - x)

input_size = X_train.shape[1]
```

```

hidden_size = 5

output_size = y_train.shape[1]

learning_rate = 0.1

epochs = 10000

np.random.seed(42)

W1 = np.random.randn(input_size, hidden_size)

B1 = np.zeros((1, hidden_size))

W2 = np.random.randn(hidden_size, output_size)

B2 = np.zeros((1, output_size))

losses = []

for epoch in range(epochs):

    Z1 = np.dot(X_train, W1) + B1

    A1 = sigmoid(Z1)

    Z2 = np.dot(A1, W2) + B2

    A2 = sigmoid(Z2)

    loss = np.mean((A2 - y_train) ** 2)

    losses.append(loss)

    dA2 = A2 - y_train

    dZ2 = dA2 * sigmoid_derivative(A2)

    dW2 = np.dot(A1.T, dZ2)

    dB2 = np.sum(dZ2, axis=0, keepdims=True)

    dA1 = np.dot(dZ2, W2.T)

    dZ1 = dA1 * sigmoid_derivative(A1)

    dW1 = np.dot(X_train.T, dZ1)

    dB1 = np.sum(dZ1, axis=0, keepdims=True)

```

```

W2 -= learning_rate * dW2

B2 -= learning_rate * dB2

W1 -= learning_rate * dW1

B1 -= learning_rate * dB1

if epoch % 1000 == 0:

    print(f"Epoch {epoch}, Loss: {loss:.4f} ")

plt.plot(losses)

plt.xlabel("Epochs")

plt.ylabel("Loss")

plt.title("Training Loss Curve")

plt.show()

def predict(X):

    Z1 = np.dot(X, W1) + B1

    A1 = sigmoid(Z1)

    Z2 = np.dot(A1, W2) + B2

    A2 = sigmoid(Z2)

    return np.argmax(A2, axis=1)

predictions = predict(X_test)

y_test_labels = np.argmax(y_test, axis=1)


accuracy = np.mean(predictions == y_test_labels)

print(f"Test Accuracy: {accuracy:.2f} ")

```

## Output

Epoch 0, Loss: 0.3792  
Epoch 1000, Loss: 0.0149  
Epoch 2000, Loss: 0.0140  
Epoch 3000, Loss: 0.0130  
Epoch 4000, Loss: 0.0121  
Epoch 5000, Loss: 0.0112  
Epoch 6000, Loss: 0.0098  
Epoch 7000, Loss: 0.0085  
Epoch 8000, Loss: 0.0084  
Epoch 9000, Loss: 0.0082



Test Accuracy: 0.97

## **Program 9**

Implement a Q-learning algorithm to navigate a simple grid environment, defining the reward structure and analyzing agent performance

### **Algorithm:**

- Define the environment (states, actions, rewards).
- Initialize Q-table with zeros.
- For each episode:
- Select an action using an epsilon-greedy policy.
- Take the action, observe reward, and update the Q-value:

$$Q(s,a)=Q(s,a)+\alpha[r+\gamma\max_{a'}Q(s',a')-Q(s,a)]$$

- Repeat until convergence and analyze agent performance.

### **Code:**

```
import numpy as np
import random
import matplotlib.pyplot as plt

GRID_SIZE = 5

ACTIONS = ['up', 'down', 'left', 'right']

ACTION_MAP = {0: (-1, 0), 1: (1, 0), 2: (0, -1), 3: (0, 1)}

GOAL_STATE = (4, 4)

PENALTY_STATE = (2, 2)

GAMMA = 0.9 # Discount factor

ALPHA = 0.1 # Learning rate

EPSILON = 0.1 # Exploration rate
```

```

EPISODES = 1000

Q_table = np.zeros((GRID_SIZE, GRID_SIZE, len(ACTIONS)))

episode_rewards = []

episode_steps = []

def take_action(state, action):

    new_state = (max(0, min(GRID_SIZE - 1, state[0] + ACTION_MAP[action][0])),
                 max(0, min(GRID_SIZE - 1, state[1] + ACTION_MAP[action][1])))

    if new_state == GOAL_STATE:

        return new_state, 10 # Reward for reaching goal

    elif new_state == PENALTY_STATE:

        return new_state, -10 # Penalty state

    else:

        return new_state, -1 # Small penalty for each move

for episode in range(EPISODES):

    state = (0, 0) # Start at top-left corner

    done = False

    total_reward = 0

    steps = 0

    while not done:

        # Choose action ( $\epsilon$ -greedy policy)

        if random.uniform(0, 1) < EPSILON:

            action = random.randint(0, len(ACTIONS) - 1) # Explore

        else:

            action = np.argmax(Q_table[state[0], state[1], :]) # Exploit

        new_state, reward = take_action(state, action)

```

```

    Q_table[state[0], state[1], action] += ALPHA * (
        reward + GAMMA * np.max(Q_table[new_state[0], new_state[1], :]) - Q_table[state[0],
state[1], action]
    )

    state = new_state

    total_reward += reward

    steps += 1

    if state == GOAL_STATE or state == PENALTY_STATE:

        done = True

    episode_rewards.append(total_reward)

    episode_steps.append(steps)

policy = np.full((GRID_SIZE, GRID_SIZE), 'X')

for i in range(GRID_SIZE):

    for j in range(GRID_SIZE):

        if (i, j) == GOAL_STATE:

            policy[i, j] = 'G'

        elif (i, j) == PENALTY_STATE:

            policy[i, j] = 'P'

        else:

            best_action = np.argmax(Q_table[i, j, :])

            policy[i, j] = ACTIONS[best_action][0].upper()

print("Optimal Policy:")

print(policy)

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)

plt.plot(episode_rewards)

```



```

plt.xlabel("Episodes")
plt.ylabel("Total Reward")
plt.title("Episode Rewards Over Time")

plt.subplot(1, 2, 2)
plt.plot(episode_steps)
plt.xlabel("Episodes")
plt.ylabel("Steps Taken")
plt.title("Steps per Episode Over Time")

```



## Output

Optimal Policy:

```

[[ 'R' 'D' 'D' 'D' 'D' ]
 [ 'R' 'R' 'R' 'R' 'D' ]
 [ 'R' 'D' 'P' 'R' 'D' ]
 [ 'R' 'D' 'D' 'D' 'D' ]
 [ 'R' 'R' 'R' 'R' 'G' ]]

```



## **Program 10**

Write a python program

- a. to perform tokenization by word and sentence using nltk.
- b. to eliminate stop words using nltk.
- c. to perform stemming using nltk.
- d. to perform Parts of Speech tagging using nltk.

### **Algorithm:**

#### **a. Tokenization (Splitting text into words/sentences)**

- Use `nltk.word_tokenize(text)` for word tokenization.
- Use `nltk.sent_tokenize(text)` for sentence tokenization.

#### **b. Remove Stop Words**

- Use `nltk.corpus.stopwords` to filter out common words like "the", "is", etc.

#### **c. Stemming (Reducing words to root form)**

- Use `nltk.stem.PorterStemmer` to transform words (e.g., "running" → "run").

#### **d. POS Tagging (Assigning part-of-speech labels to words)**

- Use `nltk.pos_tag(tokens)` to label words as nouns, verbs, etc.

### **Code:**

```
user_input = input("Enter some text: ")
upper_input = user_input.upper()
upper_input
import nltk
from nltk.stem import PorterStemmer
from nltk.tokenize import word_tokenize
text=input("Enter some text: ")
```

```

text = text.lower()

words = word_tokenize(text)

ps = PorterStemmer()

stemmed_words = [ps.stem(w) for w in words]

stemmed_words

import nltk

from nltk.tokenize import word_tokenize

text=input("Enter some text: ")

text = text.lower()

words = word_tokenize(text)

pos_tags = nltk.pos_tag(words)

pos_tags

import spacy

nlp=spacy.load('en_core_web_sm')

doc=nlp('she saw a bear')

for word in doc:

    print (word.text,word.pos_)

```

# Output

## A.

Enter some text: Once upon a time, in a quaint village nestled between rolling hills and lush forests, there lived a young girl named Aria. Aria had a special gift: she could communicate with animals. Every morning, she would wander into the forest, where the birds would sing her songs, and the deer would share their secrets. One day, a mysterious creature appeared, and Aria's life changed forever...

```
['onc',  
'upon',  
'a',  
'time',  
,,  
'in',  
'a',  
'quaint',  
'villag',  
'nestl',  
'between',
```

## B.

Enter some text: Once upon a time, in a quaint village nestled between rolling hills and lush forests, there lived a young girl named Aria. Aria had a special gift: she could communicate with animals. Every morning, she would wander into the forest, where the birds would sing her songs, and the deer would share their secrets. One day, a mysterious creature appeared, and Aria's life changed forever...

```
[('once', 'RB'),  
( 'upon', 'IN'),  
( 'a', 'DT'),  
( 'time', 'NN'),  
( ',', ','),  
( 'in', 'IN'),  
( 'a', 'DT'),  
( 'quaint', 'NN'),  
( 'village', 'NN'),  
( 'nestled', 'VBD'),  
( 'between', 'IN'),  
( 'rolling', 'VBG'),  
( 'hills', 'NNS'),  
( 'and', 'CC'),  
( 'lush', 'JJ'),  
( 'forests', 'NNS'),  
( ',', ','),  
( 'there', 'EX'),  
( 'lived', 'VBD'),  
( 'a', 'DT'),  
( 'young', 'JJ'),  
( 'girl', 'NN'),  
( 'named', 'VBN'),  
( 'aria', 'NN'),  
( '.', '.'),  
( 'aria', 'NN'),  
( 'had', 'VBD'),  
( 'a', 'DT'),  
( 'special', 'JJ'),
```

C.

---

she PRON  
saw VERB  
a DET  
bear NOUN